

An Application: Numerical Integration

NOW THAT WE KNOW HOW TO SEND MESSAGES with MPI, let's write a program that uses message passing to solve a problem: calculate a definite integral with the trapezoidal rule. If you remember the trapezoidal rule, you can skip section 4.1.

4.1 The Trapezoidal Rule

Recall that the definite integral from a to b of a nonnegative function $f(x)$ can be thought of as the area bounded by the x -axis, the vertical lines $x = a$ and $x = b$, and the graph of the function $f(x)$. See Figure 4.1.

One approach to estimating this area or integral is to partition the region into regular geometric shapes and then add the areas of the shapes. In the trapezoidal rule, the regular geometric shapes are trapezoids; each trapezoid has its base on the x -axis, vertical sides, and its top edge joining two points on the graph of $f(x)$. See Figure 4.2.

For our purposes, we'll choose all the bases to have the same length. So if there are n trapezoids, the base of each will be $h = (b - a)/n$. The base of the leftmost trapezoid will be the interval $[a, a + h]$; the base of the next trapezoid will be $[a + h, a + 2h]$; the next, $[a + 2h, a + 3h]$; etc. In general, the base of the i th trapezoid will be $[a + (i - 1)h, a + ih]$, $i = 1, \dots, n$. In order to simplify notation, let x_i denote $a + ih$, $i = 0, \dots, n$. Then the length of the left side of the i th trapezoid will be $f(x_{i-1})$, and its right side will be $f(x_i)$. See Figure 4.3. Thus, the area of the i th trapezoid will be




Figure 4.1 Definite integral of a nonnegative function




Figure 4.2 Trapezoids approximating definite integral




Figure 4.3 The i th trapezoid

$$\frac{1}{2}h[f(x_{i-1}) + f(x_i)],$$

and the area of our entire approximation will be the sum of the areas of the trapezoids:

$$\begin{aligned} & \frac{1}{2}h[f(x_0) + f(x_1)] + \frac{1}{2}h[f(x_1) + f(x_2)] + \cdots + \frac{1}{2}h[f(x_{n-1}) + f(x_n)] \\ &= \frac{h}{2}[f(x_0) + 2f(x_1) + 2f(x_2) + \cdots + f(x_n)] \\ &= [f(x_0)/2 + f(x_n)/2 + f(x_1) + f(x_2) + \cdots + f(x_{n-1})]h. \end{aligned}$$

So by putting $f(x)$ into a subprogram, we can write a *serial* program for calculating an integral using the trapezoidal rule.

```
/* Calculate definite integral using trapezoidal rule.
 * The function f(x) is hardwired.
 * Input: a, b, n.
 * Output: estimate of integral from a to b of f(x)
 *        using n trapezoids.
 */
```

```
#include <stdio.h>
```

```
main() {
    float integral; /* Store result in integral */
    float a, b; /* Left and right endpoints */
    int n; /* Number of trapezoids */
    float h; /* Trapezoid base width */
    float x;
    int i;

    float f(float x); /* Function we're integrating */

    printf("Enter a, b, and n\n");
    scanf("%f %f %d", &a, &b, &n);

    h = (b-a)/n;
    integral = (f(a) + f(b))/2.0;
    x = a;
    for (i = 1; i <= n-1; i++) {
        x = x + h;
        integral = integral + f(x);
    }
    integral = integral*h;

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %f\n",
        a, b, integral);
} /* main */
```

Table 4.1 Assignment of subintervals to processes

Process	Interval
0	$[a, a + \frac{n}{p}h]$
1	$[a + \frac{n}{p}h, a + 2\frac{n}{p}h]$
\vdots	\vdots
i	$[a + i\frac{n}{p}h, a + (i + 1)\frac{n}{p}h]$
\vdots	\vdots
$p - 1$	$[a + (p - 1)\frac{n}{p}h, b]$

```

float f(float x) {
    float return_val;
    /* Calculate f(x).  Store calculation in return_val. */
    :
    return return_val;
} /* f */

```

4.2 Parallelizing the Trapezoidal Rule

As we saw in Chapter 2, there are several approaches to parallelizing a serial program. Perhaps the simplest approach distributes the data among the processes, and each process runs essentially the same program on its share of the data. In our case, the data is just the interval $[a, b]$ and the number of trapezoids n . So we can parallelize the trapezoidal rule program by assigning a subinterval of $[a, b]$ to each process, and having that process estimate the integral of f over the subinterval. In order to calculate the integral over $[a, b]$, the processes' local calculations are added.

An obvious question here is, How does each process know which subinterval it should integrate over, and how many trapezoids it should use? In order to answer this, suppose there are p processes and n trapezoids, and, in order to simplify the discussion, also suppose that n is evenly divisible by p . Then it is natural for the first process to calculate the area of the first n/p trapezoids, the second process to calculate the area of the next n/p , etc. Recall that MPI identifies each process by a nonnegative integer. So if there are p processes, the first is process 0, the second process 1, \dots , and the last process $p - 1$. Using the notation we developed in our discussion of the serial program, we have each process calculating integrals over the subintervals indicated in Table 4.1.

Thus each process needs the following information:

- The number of processes, p
- Its rank
- The entire interval of integration, $[a, b]$
- The number of subintervals, n

Recall from Chapter 3 that the first two items can be found by calling the MPI functions `MPI_Comm_size` and `MPI_Comm_rank`. The last two items should probably be input by the user. But this (perhaps surprisingly) can raise some difficult problems. So for our first attempt at calculating the integral, let's "hardwire" these values by simply setting their values with assignment statements.

A second obvious question is, How are the individual processes' calculations added up? One straightforward approach would be to send each process's result to, say, process 0, and have process 0 do the final addition.

With these assumptions we can write our first "real" MPI program.

```
/* Parallel Trapezoidal Rule
 *
 * Input: None.
 * Output: Estimate of the integral from a to b of f(x)
 *         using the trapezoidal rule and n trapezoids.
 *
 * Algorithm:
 *   1. Each process calculates "its" interval of
 *      integration.
 *   2. Each process estimates the integral of f(x)
 *      over its interval using the trapezoidal rule.
 *   3a. Each process != 0 sends its integral to 0.
 *   3b. Process 0 sums the calculations received from
 *       the individual processes and prints the result.
 *
 * Note: f(x), a, b, and n are all hardwired.
 */
#include <stdio.h>

/* We'll be using MPI routines, definitions, etc. */
#include "mpi.h"

main(int argc, char** argv) {
    int    my_rank; /* My process rank */
    int    p;       /* The number of processes */
    float  a = 0.0; /* Left endpoint */
    float  b = 1.0; /* Right endpoint */
    int    n = 1024; /* Number of trapezoids */
    float  h;       /* Trapezoid base length */
    float  local_a; /* Left endpoint my process */
```

```

float      local_b;   /* Right endpoint my process */
int        local_n;   /* Number of trapezoids for */
                        /* my calculation */
float      integral;  /* Integral over my interval */
float      total;     /* Total integral */
int        source;    /* Process sending integral */
int        dest = 0;  /* All messages go to 0 */
int        tag = 0;
MPI_Status status;

float Trap(float local_a, float local_b, int local_n,
           float h);   /* Calculate local integral */

/* Let the system do what it needs to start up MPI */
MPI_Init(&argc, &argv);

/* Get my process rank */
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

/* Find out how many processes are being used */
MPI_Comm_size(MPI_COMM_WORLD, &p);

h = (b-a)/n;   /* h is the same for all processes */
local_n = n/p; /* So is the number of trapezoids */

/* Length of each process's interval of
 * integration = local_n*h. So my interval
 * starts at: */
local_a = a + my_rank*local_n*h;
local_b = local_a + local_n*h;
integral = Trap(local_a, local_b, local_n, h);

/* Add up the integrals calculated by each process */
if (my_rank == 0) {
    total = integral;
    for (source = 1; source < p; source++) {
        MPI_Recv(&integral, 1, MPI_FLOAT, source, tag,
                 MPI_COMM_WORLD, &status);
        total = total + integral;
    }
} else {
    MPI_Send(&integral, 1, MPI_FLOAT, dest,
             tag, MPI_COMM_WORLD);
}

/* Print the result */
if (my_rank == 0) {
    printf("With n = %d trapezoids, our estimate\n",
           n);
}

```

```

        printf("of the integral from %f to %f = %f\n",
               a, b, total);
    }

    /* Shut down MPI */
    MPI_Finalize();
} /* main */

float Trap(
    float local_a /* in */,
    float local_b /* in */,
    int local_n /* in */,
    float h /* in */) {

    float integral; /* Store result in integral */
    float x;
    int i;

    float f(float x); /* function we're integrating */

    integral = (f(local_a) + f(local_b))/2.0;
    x = local_a;
    for (i = 1; i <= local_n-1; i++) {
        x = x + h;
        integral = integral + f(x);
    }
    integral = integral*h;
    return integral;
} /* Trap */

float f(float x) {
    float return_val;
    /* Calculate f(x). */
    /* Store calculation in return_val. */
    :
    return return_val;
} /* f */

```

Observe that this program also uses the SPMD paradigm. Even though process 0 executes an essentially different set of commands from the remaining processes, it still runs the same program. The different commands are executed by branching based on the process rank.

Also note that we were careful to distinguish between variables whose contents were significant on all the processes, and variables whose contents

were only significant on individual processes. Examples of the former are `a`, `b`, and `n`. Examples of the latter are `local_a`, `local_b`, and `local_n`. Variables whose contents are significant on all the processes are sometimes called **global variables**, and variables whose contents are significant only on individual processes are sometimes called **local variables**. If you learned to program in Pascal, this terminology may at first seem somewhat confusing. However, it's usually very easy to tell from the context which meaning is implied.

It's extremely important that we, as programmers, distinguish between global and local variables: it can be very difficult or impossible to decipher a program that makes no distinction between the two. One of the most insidious things a parallel programmer can do is to use the same variable for both global and local storage with no documentation. In general, separate variables should be allocated for global and local scalar variables. For composite variables it may be necessary to use the same storage for both global and local variables. However, if this is done, it should be clearly documented.

4.3 I/O on Parallel Systems

One obvious problem with our program is its lack of generality. The function, $f(x)$, and the input data, a , b , and n , are hardwired. So if we want to change any of these, we must edit and recompile the program. Different functions can be incorporated by revising the `Trap` function so that it takes an additional parameter—a pointer to a function. Since this has nothing to do with parallel computing, we'll leave it as an exercise to modify the program to use function pointers. However, the issue of changing the input data has everything to do with parallel computing. So we should take a look at it.

In our greetings and serial trapezoidal programs we assumed that process 0 could both read from standard input (the keyboard) and write to standard output (the terminal screen). Many parallel systems provide this much I/O. In fact, many parallel systems allow all processors to both read from standard input and write to standard output. So what's the problem?

In the first place, we were careful to say “many” (not “all”) systems provide this much I/O. But even if we could say “all,” there would still be issues that need to be resolved.

Let's look at an example. Suppose we modify the trapezoidal program so that each process attempts to read the values a , b , and n by adding the statement

```
scanf("%f %f %d", &a, &b, &n);
```

Suppose also that we run the program with two processes and the user types in

```
0 1 1024
```


What happens? Do both processes get the data? Does only one? Or, even worse, does, say, process 0 get the 0 and 1, while process 1 gets the 1024?

If all the processes get the data, what happens when we write a program in which we want process 0 to read the first input value, process 1 to read the second, etc.? If only one process gets the data, what happens to the others? Is it even reasonable to have multiple processes reading data from a single terminal?

Further, what happens if several processes attempt to simultaneously write data to the terminal screen? Does the data from process 0 get printed first, then the data from process 1, etc.? Or does the data appear in some random order? Or, even worse, does the data from the different processes get all mixed up—say, half a line from 0, two characters from 1, three characters from 0, two lines from 2, etc.?

Regardless of how you feel these questions should be answered, there is not (yet) a consensus in the parallel computing world.

Thus far, we have assumed that process 0 can at least write to standard output. We will also assume that it can read from standard input. In most cases, we will only assume that process 0 can do I/O. It should be noted that this is a fairly weak assumption, since, as we noted, many parallel systems allow multiple processes to carry out I/O.¹ You might want to ask your local expert whether there are any restrictions on which processes can do I/O.

OK. If only process 0 can do I/O, then we need for process 0 to send the user input to the other processes. This is readily accomplished with a short I/O function that uses `MPI_Send` and `MPI_Recv`.

```
/* Function Get_data
 * Reads in the user input a, b, and n.
 * Input parameters:
 *   1. int my_rank: rank of current process.
 *   2. int p: number of processes.
 * Output parameters:
 *   1. float* a_ptr: pointer to left endpoint a.
 *   2. float* b_ptr: pointer to right endpoint b.
 *   3. int* n_ptr: pointer to number of trapezoids.
 * Algorithm:
 *   1. Process 0 prompts user for input and
 *      reads in the values.
 *   2. Process 0 sends input values to other
 *      processes.
 */
void Get_data(
    float* a_ptr /* out */,
    float* b_ptr /* out */,
```

¹ If your system won't allow you to do this, skip forward to Chapter 8 for a discussion of how to deal with systems with limited I/O capabilities

```

    int*    n_ptr    /* out */,
    int     my_rank  /* in  */,
    int     p        /* in  */) {

    int source = 0;    /* All local variables used by */
    int dest;         /* MPI_Send and MPI_Recv      */
    int tag;
    MPI_Status status;

    if (my_rank == 0){
        printf("Enter a, b, and n\n");
        scanf("%f %f %d", a_ptr, b_ptr, n_ptr);
        for (dest = 1; dest < p; dest++){
            tag = 0;
            MPI_Send(a_ptr, 1, MPI_FLOAT, dest, tag,
                     MPI_COMM_WORLD);
            tag = 1;
            MPI_Send(b_ptr, 1, MPI_FLOAT, dest, tag,
                     MPI_COMM_WORLD);
            tag = 2;
            MPI_Send(n_ptr, 1, MPI_INT, dest, tag,
                     MPI_COMM_WORLD);
        }
    } else {
        tag = 0;
        MPI_Recv(a_ptr, 1, MPI_FLOAT, source, tag,
                 MPI_COMM_WORLD, &status);
        tag = 1;
        MPI_Recv(b_ptr, 1, MPI_FLOAT, source, tag,
                 MPI_COMM_WORLD, &status);
        tag = 2;
        MPI_Recv(n_ptr, 1, MPI_INT, source, tag,
                 MPI_COMM_WORLD, &status);
    }
} /* Get_data */

```

Note that we used different tags for the messages containing `a_ptr`, `b_ptr`, and `n_ptr`. Although MPI guarantees that a sequence of messages sent from one process to another process will be received in the order they will be sent, this certainly doesn't hurt, and it provides a little extra assurance when we're developing the program that our messages are being received where they should be.

4.4 Summary

We didn't learn any new MPI commands in this chapter, but we did write a program that solved a problem. The process we used is typical when writing a parallel program:

1. First we recalled a serial algorithm for solving our problem: we studied the trapezoidal rule for estimating a definite integral.
2. In order to parallelize the serial algorithm, we simply partitioned the data among the processes, and each process ran essentially the same program on its data: we partitioned the interval of integration, $[a, b]$, among the processes, and each process estimated an integral over its subinterval.
3. The “local” calculations produced by the individual processes were combined to produce the final result: each process sent its integral to process 0, which summed them and printed the result.

We noted the importance of distinguishing between global variables—variables whose contents have significance on all processes—and local variables—variables whose contents only have significance on individual processes.

We briefly discussed I/O on parallel systems. We noted that many parallel systems allow each process to read from standard input and to write to standard output. However we saw that there may be difficulties in arranging that input is read by the correct process and whether output appears in the correct order. We will avoid this problem (for the time being) by only assuming that process 0 can do I/O.

4.5 References

A discussion of the trapezoidal rule and other methods of numerical integration can be found in [19]. A discussion of the issues involved in parallel I/O can be found in [9, 10].

4.6 Exercises

1. Type in the first version of the parallel trapezoidal rule program. Define $f(x)$ to be a function whose integral you can easily calculate by hand (e.g., $f(x) = x^2$). Compile it and run it with different numbers of processes. What happens if you try to run it with just one process?
2. Modify the parallel trapezoidal rule program so that a , b , and n are read in and distributed by process 0—use the `Get_data` function. Where should the function be called? Do you need to make any modifications other than including the definition of `Get_data` and a call to `Get_data`?

4.7 Programming Assignments

1. Modify the parallel trapezoidal rule program so that it has several different functions it can integrate, and the chosen function is passed to the Trap subroutine. Have the user select which function is to be integrated by giving him a menu of possible functions.
2. A more accurate alternative to the trapezoidal rule is Simpson's rule. The basic idea is to approximate the graph of $f(x)$ by arcs of parabolas rather than line segments. Suppose that $p < q$ are real numbers, and let r be the midpoint of the segment $[p, q]$. If we let $h = (q - p)/2$, then an equation for the parabola passing through the points $(p, f(p))$, $(r, f(r))$, and $(q, f(q))$ is

$$y = \frac{f(p)}{2h^2}(x - r)(x - q) - \frac{f(r)}{h^2}(x - p)(x - q) + \frac{f(q)}{2h^2}(x - p)(x - r).$$

If we integrate this from p to q , we get

$$\frac{h}{3}[f(p) + 4f(r) + f(q)].$$

Thus, if we use the same notation that we used in our discussion of the trapezoidal rule and we assume that n , the number of subintervals of $[a, b]$, is even, we can approximate

$$\begin{aligned} \int_a^b f(x)dx \doteq \frac{h}{3}[f(x_0) + 4f(x_1) + 2f(x_2) + 4f(x_3) \\ + \cdots + 2f(x_{n-2}) + 4f(x_{n-1}) + f(x_n)]. \end{aligned}$$

Assuming that n/p is even, write

- a. a serial program and
- b. a parallel program that uses Simpson's rule to estimate $\int_a^b f(x)dx$.